

Evidence-Gated Medical LLM Alerts for Pharmacogenomic Claims

Anonymous Author(s)

Abstract

LLM-generated medication alerts can sound more certain than their evidence permits. We built a deterministic pre-display monitor for one workflow: pharmacogenomic and genomic medication-alert text. The monitor checks source support, population fit, and claim strength, then returns allow, narrow, abstain, or deny. Stage 1 is a synthetic specification test: the case bank supplies the structured annotations a future live system would need to extract, and ablation tests whether any check is dispensable. Disabling source support, population fit, or claim strength let 2, 1, and 6 of 23 designed overclaims pass unchanged, respectively. With all checks active, 23/23 designed overclaim archetypes were routed and 10/10 bounded alerts passed unchanged. GPT-5.6-terra and Grok-4.5 self-evaluation baselines also routed 23/23 overclaims, but differed from the deterministic monitor on selected actions or surfaced rationales. The first-triggering check matched the author-assigned label in 31/33 cases; we interpret the two nonmatches as precedence sensitivity rather than action failure. These are specification-conformance results, not clinical safety, generalization, or patient-benefit evidence.

Keywords: clinical decision support; large language models; pharmacogenomics; evidence overclaiming; model evaluation

Institutional Review Board (IRB) Stage 1 uses only author-designed synthetic cases and public guideline/database names, and therefore does not require IRB approval. Any Stage 2 use of clinical cases or expert adjudication would require appropriate local governance.

1. Introduction

Healthcare LLMs can make evidentiary overreach look syntactically clean. A medication alert may cite a real guideline while turning a bounded pharmacogenomic recommendation into a universal instruction, a population-specific association into a claim for a different population, or a surrogate endpoint into patient-outcome benefit. Here, transportability means applying evidence from one source population as if it supported the same claim in a different target population. The gap matters because fluency can hide what evidence does, and does not, authorize. Pharmacogenomic medication alerts are a useful testbed because CPIC-style resources distinguish genotype, phenotype, evidence level, and prescribing recommendation ([Clinical Pharmacogenetics Implementation Consortium, 2026](#)). The thesis tested here is deliberately narrow and disputable: overclaim control in this workflow needs three separable checks—source support, population fit, and claim strength—rather than a single endpoint-evidence rule.

2. Evidence Boundary

Healthcare AI evaluation should distinguish implementation consistency from patient-outcome benefit. Byrne et al. emphasize that outcome claims require pragmatic clinical evaluation ([Byrne et al., 2024](#)), while Stegenga motivates caution when evidentiary pipelines make interventions appear stronger than they

Data and Code Availability The Stage 1 case bank is synthetic and contains no real patient records, protected health information, or private genomic data. For anonymous venue review, code and synthetic data should be supplied as an anonymized supplemental archive; the de-anonymized public repository may be disclosed after acceptance or in course-facing materials.

are (Stegenga, 2018). FDA-NIH BEST materials provide endpoint vocabulary: biomarkers, surrogate endpoints, validated surrogates, and clinical outcomes are not interchangeable (U.S. Food and Drug Administration and National Institutes of Health, 2016). The three checks used here instantiate a more general evidence-boundary framework for clinical LLM alerts: source support asks whether cited evidence exists and applies, population fit asks whether the evidence transports to the target patient or population, and claim strength asks whether the evidence supports the requested action. Pharmacogenomics supplies concrete examples rather than a one-off exception: CPIC-style resources can support bounded prescribing language, including CYP2C19-clopidogrel and DPYD-fluoropyrimidine examples (Lee et al., 2022; Amstutz et al., 2018), but do not automatically support claims that an LLM improves outcomes, that an association is causal, or that a result generalizes across populations without recalibration.

3. Related Work

LLM factuality methods such as SelfCheckGPT test whether generated content is consistent across black-box samples (Manakul et al., 2023). Selective classification supplies the decision vocabulary for abstention: a system may decline to answer when coverage is less important than risk control (Geifman and El-Yaniv, 2017). Clinical alert fatigue explains why this choice matters in practice: repeated and low-information alerts can reduce clinician response. This study tests a smaller executable claim. Given structured annotations about source support, target population, and claim strength, a pre-display monitor can separate designed pharmacogenomic overclaims from bounded alerts while showing which checks an endpoint-only rule would miss.

4. Methods

The controller is a deterministic pre-display monitor, not an autonomous clinical agent. A Stage 1 case contains a candidate alert, source-status annotation, population-fit annotation, endpoint annotation, actionability annotation, and expected

review label. The monitor reads those annotations in a fixed order: source support, population fit, then claim strength. It returns one of five actions: allow, narrow, abstain (population fit), abstain (source support), or deny. The repository preserves the implementation identifiers for the three checks, but the paper uses their clinical names. Table 1 maps the general evidence-boundary pattern to the pharmacogenomic implementation; the code expands each boundary into exact predicates.

5. Stage 1 Scaffold

The case bank contains 33 author-designed synthetic cases: 10 bounded aligned alerts and 23 overclaim archetypes involving universal action language, unsupported genetics, population mismatch, citation weakness, variant uncertainty, obsolete evidence, and changed clinical context. An isolating case is designed so that only one check changes the action; with that check disabled, the overclaim passes unchanged. Within the 23 overclaim archetypes, three cases isolate a fabricated CPIC addendum attached to an otherwise actionable claim, a real guideline-supported claim applied to a mismatched population, and a stale but real guideline used as current evidence. The 23/10 split is stress-test coverage, not a clinical base rate. Each case has expected labels, but the controller does not read them. The three checks are summarized in Table 1. The decisive test is ablation, not action conformance: if a disabled check lets an overclaim pass unchanged, that check is doing work not captured by the remaining checks.

With all checks active, all 23 designed overclaim archetypes were blocked or narrowed and all 10 bounded alerts passed unchanged. Table 2 shows that disabling each check allows at least one designed overclaim to pass unchanged, so source support and population fit are not reducible to claim strength. Removing source support allows 2/23 overclaims through unchanged; removing population fit allows 1/23 through; removing claim strength allows 6/23 through. No check is redundant in this synthetic case bank: each check blocks at least one designed overclaim that passes unchanged when that check is disabled.

Table 1: Evidence-boundary framework and pharmacogenomic instantiation.

Boundary	General question	Failure mode	PGx implementation
Source support	Does the cited evidence exist, remain current, and apply to the stated context?	Hallucinated, absent, conflicting, or outdated citation.	CPIC, ClinPGx, or ClinVar anchor is verified, stale, conflicting, unsupported, or unverifiable.
Population fit	Does the evidence transport to this target patient or population?	Source-target mismatch, ancestry/founder effect overreach, or unrepresented population.	HLA, DPYD, Genome India, Israel founder-variant, or All of Us/TOPMed transport flags.
Claim strength	Does the evidence support the displayed action?	Surrogate endpoint, association, uncertain variant, or polygenic score treated as treatment authority.	CYP2C19–clopidogrel and DPYD–fluoropyrimidine alerts may be bounded; unsupported medication actions are narrowed or denied.

LLM self-evaluation baseline Two frontier LLMs classified the same 33 structured annotations without seeing the deterministic rule table. This tests a natural objection: perhaps the model can police its own evidence boundary if prompted carefully. Both baselines routed all 23 designed overclaims and allowed all 10 bounded alerts, but they were not identical to the deterministic monitor. GPT-5.6-terra matched 31/33 actions and 25/33 primary checks; Grok-4.5 matched 31/33 actions and 32/33 primary checks. Thus, on this structured synthetic task, LLM self-review did not simply fail; the remaining issue is that stochastic reviewers can choose different rationales or stronger actions without an explicit governance policy.

Precedence sensitivity as design Primary-check conformance, defined as the first enabled check matching the author-assigned check label, was 31/33. In PGX19 and PGX24, the controller returned the expected action, but the surfaced explanation differed because the current policy checks source support before population fit or claim strength. These cases are better understood as precedence-sensitive governance examples than as safety failures. For PGX24, Table 2 shows that the action may remain restrictive while the explanation changes from source support to population transport to unsupported medication action.

6. Evaluation Plan

Stage 2 turns the specification test into a held-out reader study. Independent case authors write pharmacogenomic alert scenarios without seeing the decision table. The ungated arm shows the LLM draft unchanged; the gated arm shows the same draft after structured annotations pass through the monitor. Two clinical reviewers assign binary labels for overclaim present, denial appropriate, clinician authority preserved, population fit adequate, and source support adequate; they also select one error class from a fixed taxonomy. The primary endpoint is the difference in reviewer-labeled overclaim rate between arms. The safety endpoint is the difference in inappropriate-denial rate; the monitor is useful only if overclaiming falls without a compensating rise in inappropriate denial.

7. Ethics and Limitations

Transparency The monitor produces a decision trace. For each alert, the trace names the structured evidence feature that triggered allow, narrow, abstain, or deny. Clinicians and reviewers can inspect the rule that changed, blocked, or preserved the displayed alert.

Fairness Fairness and bias are central in pharmacogenomics. HLA-B*15:02 risk language depends on ancestry and tested population,

Table 2: Stage 1 ablation, LLM self-evaluation, and precedence-sensitivity results.

Analysis	Result	Interpretation
One-check ablation	Source, population, and claim removal let 2/23, 1/23, and 6/23 overclaims pass unchanged.	Each boundary catches at least one archetype that the others do not catch.
GPT-5.6-terra baseline	23/23 overclaims routed; 10/10 bounded alerts allowed; 31/33 action agreement; 25/33 check agreement.	Frontier self-review was conservative, but rationale selection varied.
Grok-4.5 baseline	23/23 overclaims routed; 10/10 bounded alerts allowed; 31/33 action agreement; 32/33 check agreement.	A second frontier self-review reproduced overclaim routing but not every action choice.
PGX24 precedence	Source-first: deny/source; population-first: abstain/population; claim-first: deny/claim.	The same case can surface different clinician-facing explanations under different governance orders.

founder-variant evidence may not generalize to the US population, and biobank results may be unrepresentative for underrepresented or recently admixed groups. Cases PGX09, PGX10, PGX25, PGX26, and PGX27 treat population fit as an equity constraint. PGX24 represents the high-risk practice of using a polygenic score to deny opioid or stimulant therapy.

Governance and limitations Stage 1 contains no PHI and makes no diagnosis or treatment recommendation; deployed clinical decision support would still need accountability for extraction, citation verification, thresholds, overrides, and audit logging. The key limitation is that results are synthetic and author-designed. They show rule conformance, not generalization, and assume structured annotations rather than extracting them from live LLM output or EHR context.

8. Conclusion

The practical problem is annotation, not compute. Stage 1 shows that the monitor logic is auditable and that the three checks do different work on a synthetic pharmacogenomic case bank. Source support, population fit, and claim strength each block at least one designed overclaim that passes when that check is removed. The LLM baselines suggest that self-evaluation over structured annotations can be conservative, but action choice and explanation order still vary without explicit predicates. That result turns the next experiment into a concrete measurement problem: determine whether live LLM and EHR

pipelines can populate the three annotations reliably enough to reduce overclaiming without increasing inappropriate denial.

References

- Ursula Amstutz et al. Clinical pharmacogenetics implementation consortium guideline for dihydropyrimidine dehydrogenase genotype and fluoropyrimidine dosing: 2017 update. *Clinical Pharmacology & Therapeutics*, 2018.
- D. W. Byrne, H. J. Domenico, and R. P. Moore. Artificial intelligence for improved patient outcomes: The pragmatic randomized controlled trial is the secret sauce. *Korean Journal of Radiology*, 2024.
- Clinical Pharmacogenetics Implementation Consortium. Clinical Pharmacogenetics Implementation Consortium Guidelines. <https://cpicpgx.org/guidelines/>, 2026.
- Yonatan Geifman and Ran El-Yaniv. Selective Classification for Deep Neural Networks. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Craig R. Lee et al. Clinical pharmacogenetics implementation consortium guideline for CYP2C19 genotype and clopidogrel therapy: 2022 update. *Clinical Pharmacology & Therapeutics*, 2022.
- Potsawee Manakul, Adian Liusie, and Mark Gales. Self-CheckGPT: Zero-Resource Black-Box Hallucination Detection for Generative Large Language Models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 9004–9017. Association for Computational Linguistics, 2023. doi: 10.18653/v1/2023.emnlp-main.557.
- Jacob Stegenga. *Medical Nihilism*. Oxford University Press, 2018.
- U.S. Food and Drug Administration and National Institutes of Health. BEST (Biomarkers, EndpointS, and Other Tools) Resource. <https://www.ncbi.nlm.nih.gov/books/NBK326791/>, 2016.