

Summary Sheet

Pruned evidence-gated pharmacogenomic LLM-alert capstone

Author: Ravi Bajaj | Course: AI in Healthcare

Conference / Symposium

Field	Value
Name	Machine Learning for Health Symposium (ML4H 2026)
URL	https://ml4h.ahli.cc/
Author instructions / CFP	https://ml4h.ahli.cc/submit/call-for-papers/
Submission deadline	September 10, 2026, 11:59 PM AoE
Author notification	October 22, 2026
Camera-ready deadline	November 7, 2026 (tentative)
Event dates	December 6-7, 2026
Location	Sydney, Australia
Suggested track	Findings Track if kept as preliminary construct-validity work; Proceedings Track only after independent review.

Title and Abstract

Title: Evidence-Gated Medical LLM Alerts for Pharmacogenomic Claims: Detecting Overclaiming Relative to Guideline-Supported Evidence

Abstract: This capstone studies a narrowed clinical AI safety problem: LLM-drafted pharmacogenomic and genomic medication alerts can overstate what guideline-supported evidence actually permits. I implement a synthetic pre-display evidence gate with three checks: endpoint/actionability, population fit, and citation/guideline support. In 30 synthetic cases, ungated drafts overclaimed evidence in 20 cases, while the gated controller reduced remaining overclaiming to 0, allowed 10 aligned bounded alerts, and produced no inappropriate denials. These results show implementation consistency and construct validity, not clinical utility. A Stage 2 study would compare ungated and gated alerts on independently authored cases with blinded reviewer adjudication of overclaiming, inappropriate denial, calibration, and error categories.

Generative AI Use

Generative AI tools were used for brainstorming, critique, code scaffolding, drafting, packaging, and formatting. The author selected the final scope, directed the pruning, verified sources, reviewed generated text, ran local tests, and remains responsible for all claims, limitations, and submission materials.

Optional Data, Code, and Tools

- Synthetic 30-case pharmacogenomic/genomic medication-alert case bank.
- Executable three-gate controller and tests.
- CSV/JSON results for overclaim detection, inappropriate denial, calibration, and gate/error categories.
- No real patient data, no protected health information, no diagnosis, and no treatment recommendation.